

Big Data




Big Data Engineering with Hadoop & Spark

Kafka Introduction



APACHE
kafka®

A distributed streaming platform



APACHE
Spark™

Session 23: Assignment 23.1

This assignment is aimed at consolidating the concepts that was learnt during the Apache Kafka session of the course.

Dataset:

Download the dataset from this [link](#).

This is how the dataset looks:

```

1 ItemTopic-{"item_id":"101"}-{"user_id":"U101"}
2 UserTopic-{"name":"John"}-{"exp":16}
3 ItemTopic-{"item_id":"101"}-{"user_id":"U106"}
4 UserTopic-{"name":"Mark"}-{"exp":18}
5 ItemTopic-{"item_id":"102"}-{"user_id":"U110"}
6 UserTopic-{"name":"Cylin"}-{"exp":15}
7 ItemTopic-{"item_id":"102"}-{"user_id":"U101"}
8 UserTopic-{"name":"Prod"}-{"exp":14}
9 ItemTopic-{"item_id":"104"}-{"user_id":"U102"}
10 UserTopic-{"name":"Abhay"}-{"exp":17}
11 ItemTopic-{"item_id":"107"}-{"user_id":"U104"}
12 UserTopic-{"name":"Misano"}-{"exp":19}
  
```

This file has two topics namely:

1. **ItemTopic:** It has item_id and user_id
2. **UserTopic:** It has user details like name and experience

Task 1:

Create a java program MyKafkaProducer.java that takes a file name and delimiter as input arguments. It should read the content of file line by line. Fields in the file are in following order:

1. Kafka Topic Name
 2. Key
 3. value
- For every line, insert the key and value to the respective Kafka broker in a fire and forget mode.
 - After record is sent, it should print appropriate message on screen.
 - Pass dataset_producer.txt as the input file and - as delimiter.

Solution:

1. Program to perform this task is as below:

- *Imports required for the program is given below:*

```

import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.ProducerRecord;
import java.io.BufferedReader;
import java.io.FileReader;
import java.io.IOException;
import java.util.Properties;
  
```

- *This is the class called “MyKafkaProducer” which takes two arguments (Input File name and delimiter) in the command line*

```
public class MyKafkaProducerLive
{
    public static void main(String[] args) throws IOException
    {
        if (args.length != 2)
        {
            System.out.println("Please provide appropriate
command line arguments");
            System.exit(-1);
        }
    }
}
```

- *We configure the properties for KafkaProducer:*
 - *We create a new instance of Properties called props*
 - *Using this instance we add properties to kafkaProducer like, bootstrapserver/meta-data-brokerlist, key and value serializers*

```
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("key.serializer",
"org.apache.kafka.common.serialization.StringSerializer");
props.put("value.serializer",
"org.apache.kafka.common.serialization.StringSerializer");
```

- *We then instantiate the KafkaProducer class called producer, we have mentioned string in <> because both key and value are String*
- *We add the properties instance (props) to KafkaProducer instance*
- *We also instantiate ProducerRecord as producerRecord*

```
KafkaProducer<String, String> producer = new
KafkaProducer<>(props);
ProducerRecord<String, String> producerRecord = null;
```

- *Now we take the data provided in the command line i.e. file name and delimiter and save them in the array of string variables called filename and delimiter*

```
String fileName = args[0];
String delimiter = args[1];
```

- We read the contents of the input file, and save their contents arrays in different variables:
 - We save the topic name, i.e. first part of array(0th index elements) in String variable topic and similarly we save key and value variables too

```
try(BufferedReader br = new BufferedReader(new
FileReader(fileName)))
{
    for(String line; (line = br.readLine()) != null; )
    {
        String[] tempArray = line.split(delimiter);
        String topic = tempArray[0];
        String key = tempArray[1];
        String value = tempArray[2];
```

- Now, we pass the variables topic, key and value to producer record
- We also print appropriate message which shows the topics, key and value contents
- We finally, close the producer

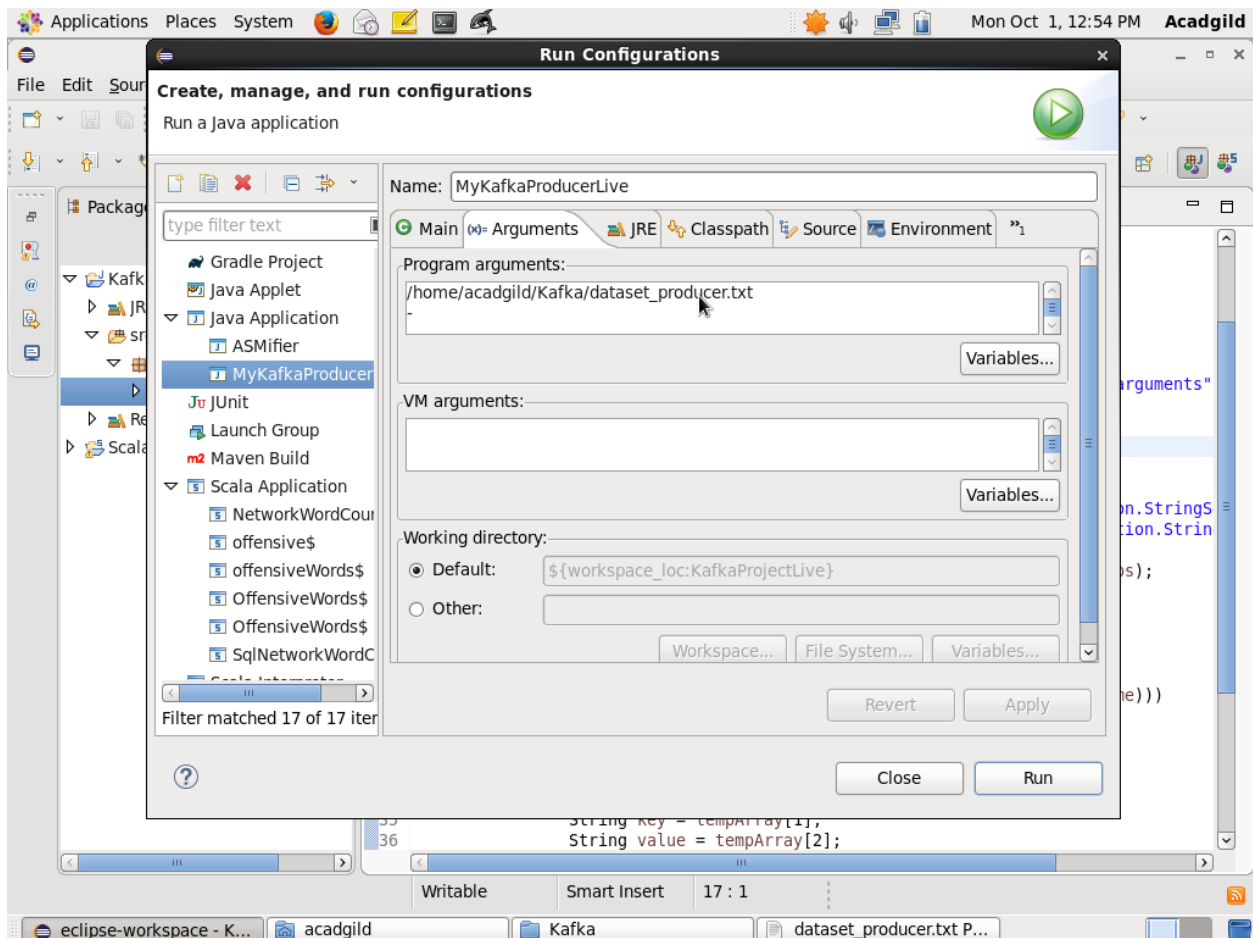
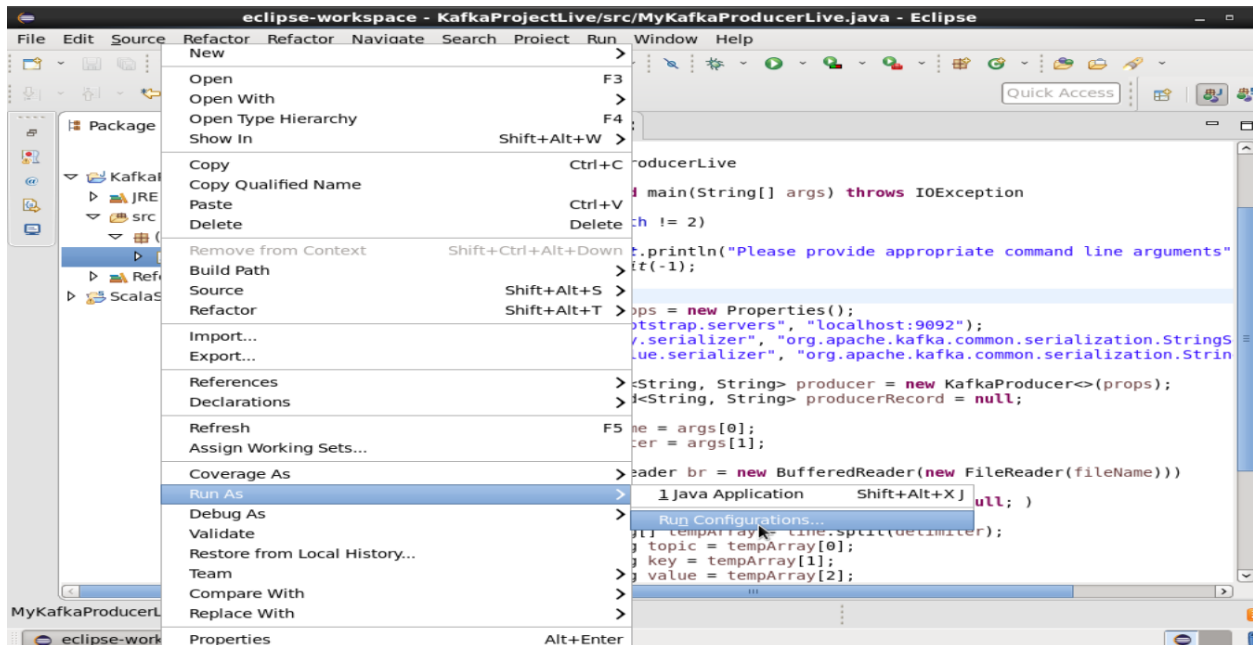
```
        producerRecord = new ProducerRecord<String,
String>(topic, key, value);
        producer.send(producerRecord);
        System.out.printf("Record sent to topic:%s.
Key:%s, Value:%s\n", topic, key, value);
    }
}

producer.close();
}
```

2. Next, we start Zookeeper and Kafka Server using the following commands:

- \$ \$KAFKA_HOME/bin/zookeeper-server-start.sh \$KAFKA_HOME/config/zookeeper.properties
- \$ \$KAFKA_HOME/bin/kafka-server-start.sh \$KAFKA_HOME/config/server.properties

3. Next, we pass the arguments in “Run Configurations”, and click on “Run” as shown below:




```

<terminated> MyKafkaProducerLive [Java Application] /usr/java/jdk1.8.0_151/bin/java (Oct 1, 2018, 12:43:44 PM)
Record sent to topic:ItemTopic. Key:{"item_id":"101"}, Value:{"user_id":"U101"}
Record sent to topic:UserTopic. Key:{"name":"John"}, Value:{"exp":16}
Record sent to topic:ItemTopic. Key:{"item_id":"101"}, Value:{"user_id":"U106"}
Record sent to topic:UserTopic. Key:{"name":"Mark"}, Value:{"exp":18}
Record sent to topic:ItemTopic. Key:{"item_id":"102"}, Value:{"user_id":"U110"}
Record sent to topic:UserTopic. Key:{"name":"Cylin"}, Value:{"exp":15}
Record sent to topic:ItemTopic. Key:{"item_id":"102"}, Value:{"user_id":"U101"}
Record sent to topic:UserTopic. Key:{"name":"Prod"}, Value:{"exp":14}
Record sent to topic:ItemTopic. Key:{"item_id":"104"}, Value:{"user_id":"U102"}
Record sent to topic:UserTopic. Key:{"name":"Abhay"}, Value:{"exp":17}
Record sent to topic:ItemTopic. Key:{"item_id":"107"}, Value:{"user_id":"U104"}
Record sent to topic:UserTopic. Key:{"name":"Misano"}, Value:{"exp":19}

```

4. Next, we run the console consumer commands on terminal to view the output of the program, using the below command:

- To read contents of ItemTopic:

```
$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic ItemTopic --from-beginning --zookeeper localhost:2181 --property print.key=true
```

- To read contents of UserTopic:

```
$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic UserTopic --from-beginning --zookeeper localhost:2181 --property print.key=true
```

```

[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic ItemTopic --from-beginning --zookeeper localhost:2181 --property print.key=true
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major release. Consider using the new consumer by passing [bootstrap-server] instead of [zookeeper].
{"item_id":"101"} {"user_id":"U101"}
{"item_id":"101"} {"user_id":"U106"}
{"item_id":"102"} {"user_id":"U110"}
{"item_id":"102"} {"user_id":"U101"}
{"item_id":"104"} {"user_id":"U102"}
{"item_id":"107"} {"user_id":"U104"}
^CProcessed a total of 6 messages
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic UserTopic --from-beginning --zookeeper localhost:2181 --property print.key=true
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major release. Consider using the new consumer by passing [bootstrap-server] instead of [zookeeper].
{"name":"John"} {"exp":16}
{"name":"Mark"} {"exp":18}
{"name":"Cylin"} {"exp":15}
{"name":"Prod"} {"exp":14}
{"name":"Abhay"} {"exp":17}
{"name":"Misano"} {"exp":19}
^CProcessed a total of 6 messages
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$

```

Task 2:

Modify the previous program `MyKafkaProducer.java` and create a new Java program `KafkaProducerWithAck.java`.

- This should perform the same task as of `KafkaProducer.java` with some modification.
 - When passing any data to a topic, it should wait for acknowledgement.
 - After acknowledgement is received from the broker, it should print the key and value which has been written to a specified topic.
 - The application should attempt for 3 retries before giving any exception.
- Pass `dataset_producer.txt` as the input file and `-as` delimiter.

Solution:

1. Program to perform the task:

- *Imports required for the program is given below:*

```
import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.ProducerRecord;
import java.io.BufferedReader;
import java.io.FileReader;
import java.io.IOException;
import java.util.Properties;
import java.util.concurrent.ExecutionException;
```

- *This is the class called “MyKafkaProducerWithAck” which takes two arguments (Input File name and delimiter) in the command line*

```
public class MyKafkaProducerWithAck
{
    public static void main(String[] args) throws IOException,
        InterruptedException, ExecutionException
    {
        if (args.length != 2)
        {
            System.out.println("Please provide appropriate
command line arguments");
            System.exit(-1);
        }
    }
}
```


- *We configure the properties for KafkaProducer:*
 - *We create a new instance of Properties called props*
 - *Using this instance we add properties to kafkaProducer like, bootstrapserver/meta-data-brokerlist, key and value serializers,acks and retries*

- *Acks “all”- this means that the producer will receive a success response from the broker once all in-sync replicas received the message*
 - *Retries 3- When the producer receives an error message from the server, the error could be transient (e.g., a lack of leader for a partition). In this case, the value of the retries parameter will control how many times the producer will retry sending the message before giving up and notifying the client of an issue*

```
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("acks", "all");
props.put("retries", 3);
props.put("key.serializer",
"org.apache.kafka.common.serialization.StringSerializer");
props.put("value.serializer",
"org.apache.kafka.common.serialization.StringSerializer");
```

- *We then instantiate the KafkaProducer class called producer, we have mentioned string in <> because both key and value are String*
- *We add the properties instance (props)to KafkaProducer instance*
- *We also instantiate ProducerRecord as producerRecord*

```
KafkaProducer<String, String> producer = new
KafkaProducer<>(props);
ProducerRecord<String, String> producerRecord = null;
```

- *Now we take the data provided in the command line i.e. file name and delimiter and save them in the array of string variables called filename and delimiter*

```
String fileName = args[0];
String delimiter = args[1];
```

- *We read the contents of the input file, and save their contents arrays in different variables:*
 - *We save the topic name i.e. first part of array(0th index elements) in String variable topic and similarly we save key and value variables too*

```
try(BufferedReader br = new BufferedReader(new
FileReader(fileName)))
{
    for(String line; (line = br.readLine()) != null; )
    {
        String[] tempArray = line.split(delimiter);
        String topic = tempArray[0];
        String key = tempArray[1];
        String value = tempArray[2];
```

- *Now, we pass the variables topic, key and value to producer record.*
- *We also print appropriate message which shows the topics, key and value contents.*
- *We finally, close the producer*

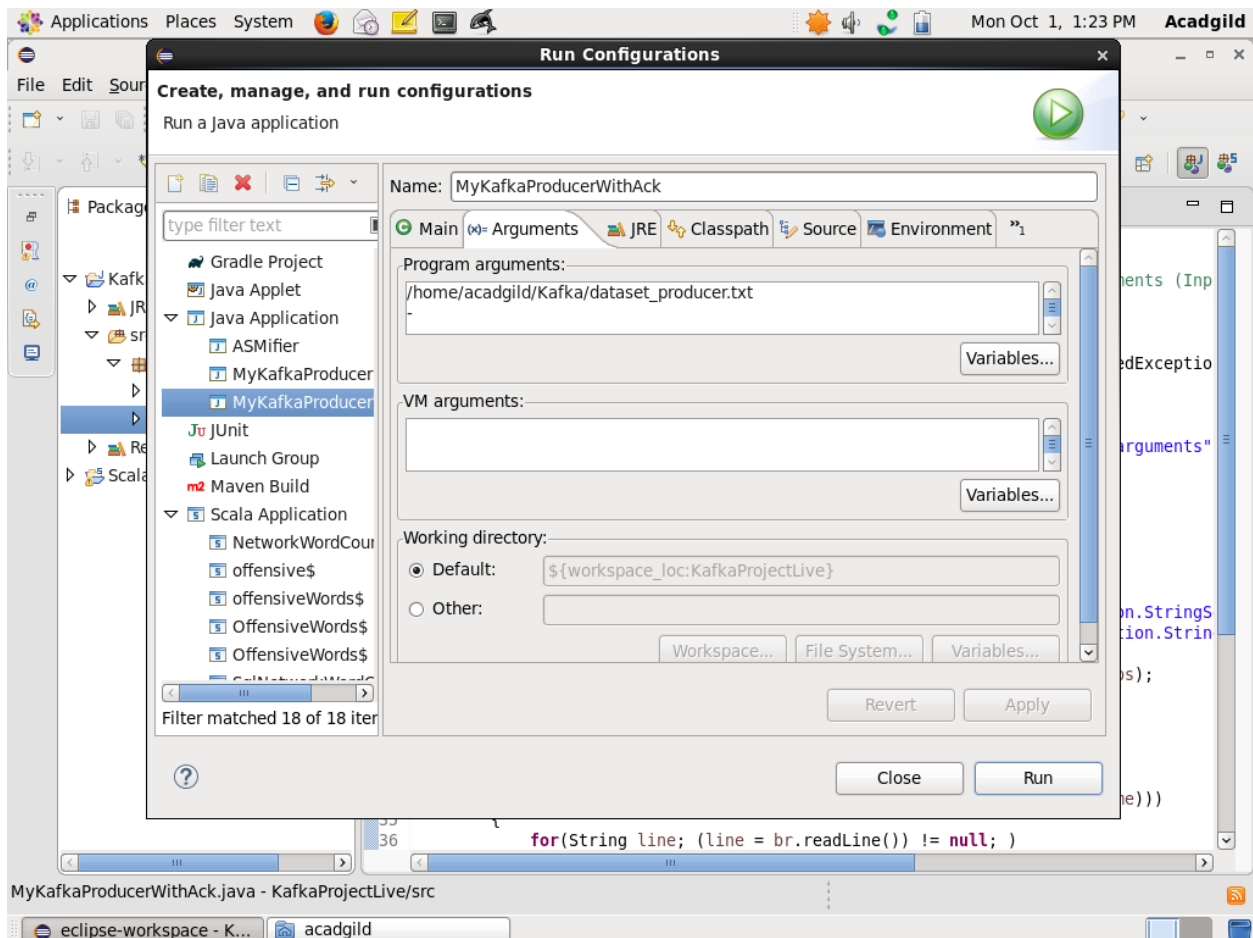
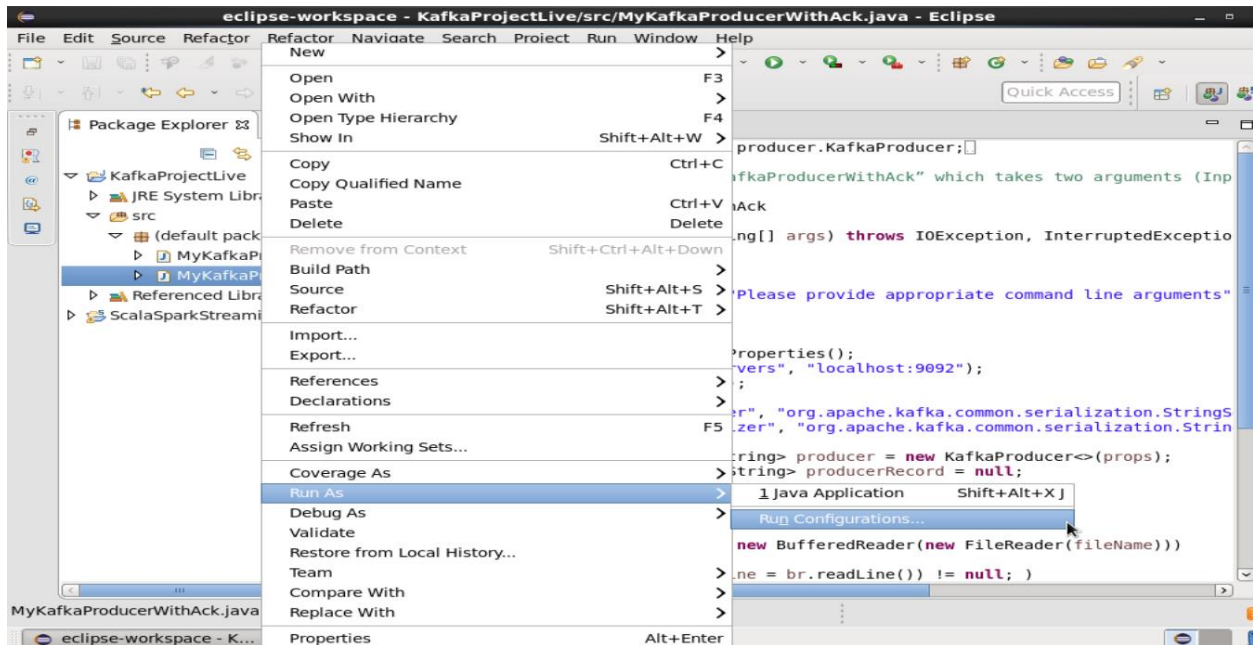
```
        producerRecord = new ProducerRecord<String,
String>(topic, key, value);
        producer.send(producerRecord).get();
        System.out.printf("Record sent to topic:%s and
acknowledged as well. Key:%s,Value:%s\n", topic, key, value);
    }
}

producer.close();
}
```

2. Next, we start Zookeeper and Kafka Server using the following commands:

- *\$ \$KAFKA_HOME/bin/zookeeper-server-start.sh \$KAFKA_HOME/config/zookeeper.properties*
- *\$ \$KAFKA_HOME/bin/kafka-server-start.sh \$KAFKA_HOME/config/server.properties*

3. Next, we pass the arguments in “Run Configurations”, and click on “Run” as shown below:



```

<terminated> MyKafkaProducerWithAck [Java Application] /usr/java/jdk1.8.0_151/bin/java (Oct 1, 2018, 1:16:47 PM)
Record sent to topic:ItemTopic and acknowledged as well. Key:{item_id:"101"},Value:{user_id:"U101"}
Record sent to topic:UserTopic and acknowledged as well. Key:{name:"John"},Value:{exp:16}
Record sent to topic:ItemTopic and acknowledged as well. Key:{item_id:"101"},Value:{user_id:"U106"}
Record sent to topic:UserTopic and acknowledged as well. Key:{name:"Mark"},Value:{exp:18}
Record sent to topic:ItemTopic and acknowledged as well. Key:{item_id:"102"},Value:{user_id:"U110"}
Record sent to topic:UserTopic and acknowledged as well. Key:{name:"Cylin"},Value:{exp:15}
Record sent to topic:ItemTopic and acknowledged as well. Key:{item_id:"102"},Value:{user_id:"U101"}
Record sent to topic:UserTopic and acknowledged as well. Key:{name:"Prod"},Value:{exp:14}
Record sent to topic:ItemTopic and acknowledged as well. Key:{item_id:"104"},Value:{user_id:"U102"}
Record sent to topic:UserTopic and acknowledged as well. Key:{name:"Abhay"},Value:{exp:17}
Record sent to topic:ItemTopic and acknowledged as well. Key:{item_id:"107"},Value:{user_id:"U104"}
Record sent to topic:UserTopic and acknowledged as well. Key:{name:"Misano"},Value:{exp:19}

```

4. Next, we run the console consumer commands on terminal to view the output of the program, using the below command:

- To read contents of ItemTopic:

```
$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic ItemTopic --from-beginning --zookeeper localhost:2181 --property print.key=true
```

```

[acadgild@localhost ~]$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic ItemTopic --from-beginning --zookeeper localhost:2181 --prop
erty print.key=true
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major release. Consider using the new consum
er by passing [bootstrap-server] instead of [zookeeper].
{"item_id":"101"} {"user_id":"U101"}
{"item_id":"101"} {"user_id":"U106"}
{"item_id":"102"} {"user_id":"U110"}
{"item_id":"102"} {"user_id":"U101"}
{"item_id":"104"} {"user_id":"U102"}
{"item_id":"107"} {"user_id":"U104"}
{"item_id":"101"} {"user_id":"U101"}
{"item_id":"101"} {"user_id":"U106"}
{"item_id":"102"} {"user_id":"U110"}
{"item_id":"102"} {"user_id":"U101"}
{"item_id":"104"} {"user_id":"U102"}
{"item_id":"107"} {"user_id":"U104"}
{"item_id":"101"} {"user_id":"U101"}
{"item_id":"101"} {"user_id":"U106"}
{"item_id":"102"} {"user_id":"U110"}
{"item_id":"102"} {"user_id":"U101"}
{"item_id":"104"} {"user_id":"U102"}
{"item_id":"107"} {"user_id":"U104"}

```

- To read contents of UserTopic:

\$ \$KAFKA_HOME/bin/kafka-console-consumer.sh --topic UserTopic --from-beginning --zookeeper localhost:2181 --property print.key=true

```
[ajitgild@localhost ~]$ $KAFKA_HOME/bin/kafka-console-consumer.sh --topic UserTopic --from-beginning --zookeeper localhost:2181 --property print.key=true
Using the ConsoleConsumer with old consumer is deprecated and will be removed in a future major release. Consider using the new consumer by passing [bootstrap-server] instead of [zookeeper].
{"name":"John"} {"exp":16}
{"name":"Mark"} {"exp":18}
{"name":"Cylin"} {"exp":15}
{"name":"Prod"} {"exp":14}
{"name":"Abhay"} {"exp":17}
{"name":"Misano"} {"exp":19}
{"name":"John"} {"exp":16}
{"name":"Mark"} {"exp":18}
{"name":"Cylin"} {"exp":15}
{"name":"Prod"} {"exp":14}
{"name":"Abhay"} {"exp":17}
{"name":"Misano"} {"exp":19}
{"name":"John"} {"exp":16}
{"name":"Mark"} {"exp":18}
{"name":"Cylin"} {"exp":15}
{"name":"Prod"} {"exp":14}
{"name":"Abhay"} {"exp":17}
{"name":"Misano"} {"exp":19}
```